



PROMETHEUS

MONITORING SISTEMA I SERVISA



Maša Nešić, I 6774

MONITORING

ZA POTREBE MIKROSERVISA I INTERNETA STVARI I SERVISA

- Šta posmatramo
 - Metriku hardvera uređaja
 - Firmware
 - Kod aplikacije
 - Spoljašnje okruženje
 - Interakciju između uređaja i cloud-a
 - Pomoćne sisteme
- Zašto posmatramo?
 - Detektovanje, debugiranje i rešavanje problema koji se javljaju u sistemima, dok ti sistemi rade
 - Dugoročna analitika i izgradnja modela
 - Optimizacija rada sistema
 - Praćenje promena za potrebe donošenja poslovnih i tehničkih odluka
 - Za potrebe drugih sistema/procesa (QA, automatizacija, bezbednost, ...)

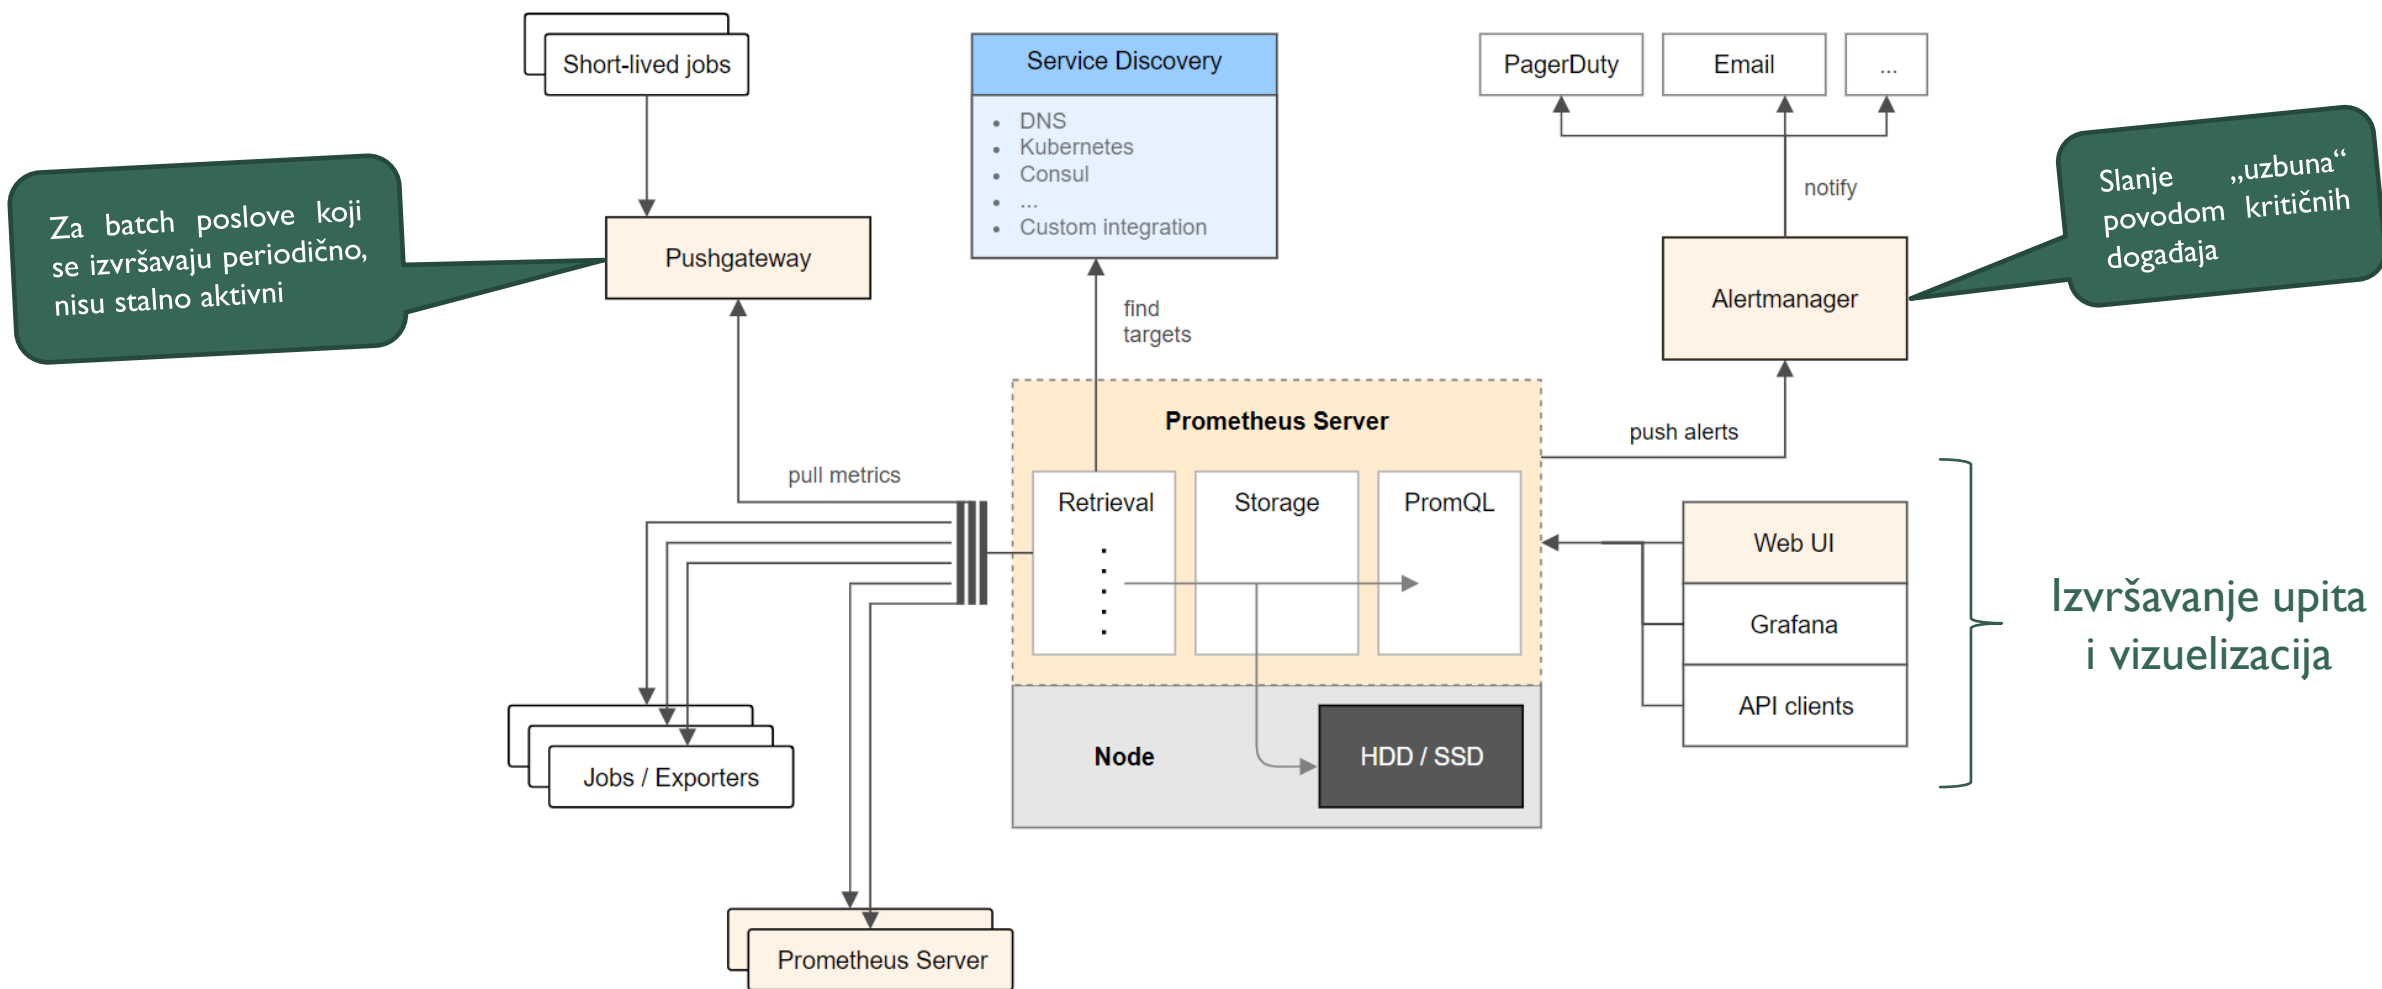
PROMETHEUS



- Sistem za monitoring sistema i servisa, i baza za smeštanje vremenskih serija podataka
- Open-source projekat
- Inicijalno razvijen od strane SoundCloud-a
- Zvanični projekat Cloud Native Computing Fondacije



ARHITEKTURA



Prometheus architecture

GLAVNE KARAKTERISTIKE


- Multi-dimenzionalni model podataka - vremenska serija je definisana imenom metrike, a labelle (key-value parovi) dodaju dimenzije
- **PromQL** – moćan i fleksibilan jezik za pisanje upita koji može da iskoriti prednosti multi-dimenzionalnosti
- HTTP pull model za prikupljanje metrika
- Podrška za monitoring „batch“ poslova preko posrednika (**Pushgateway**)
- Podrška za vizuelizaciju – kreiranje grafova, tabela i dashboard-ova
- Efikasno skladištenje – lokalna on-disk baza (nezavisno od distribuiranih skladišta, autonomne instance servera), sa mogućnošću integracije sa udaljenim sistemima za skladištenje
- Precizno kreiranje alertova uz pomoć PromQL-a
- Klijentske biblioteke
- Veliki broj raspoloživih eksportera za integraciju third-party podataka sa Prometheus-om
<https://prometheus.io/docs/instrumenting/exporters/>

SKLADIŠTENJE

- Uzorci se grupišu u blokove materijala koji je skupljan u protekla 2h
- Svaki blok se sastoji od jednog ili više chunk fajlova koji sadrže vremenske serije podataka prikupljenih u tom periodu, fajl sa metapodacima i indeksni fajl (indeksira nazive metrika i labela za chunk fajlove)
- Kada se neka serija obriše preko API-ja, informacija o tome se čuva u zasebnom fajlu (ne vrši se odmah brisanje i preuređivanje podataka iz chunk fajlova)
- Trenutni blok, u koji se smeštaju podaci se čuva u memoriji i loguje u originalnom formatu u write-ahead fajl (kao siguronosni mehanizam u slučaju otkaza)
- Vremenom se vrši kompakcija 2h blokova u blokove koji sadrže veće količine infomacija - do 10% ukupnog vremena za koje se čuvaju podaci ili 31 dan (šta god je manje)
- Moguće je podešavati koliko dugo se čuvaju podaci – default je 15 dana
- Ukupan potreban prostor = vreme čuvanja podataka (u sekundama) * broj uzoraka koji se prikuplja u sekundi * broj B po uzorku (najčešće 1-2 B)

REMOTE SKLADIŠTENJE

- Korišćenjem lokalnog skladišta, svaka instanca je ograničena na sopstvenu pouzdanost i skalabilnost
- Kao rešenje tog problema, Prometheus nudi interfejs za integraciju sa remote sistemima za skladištenje
- Obezbeđuju se upis i čitanje sa remote URL u standardizovanom formatu
- Komunikacija sa adapterom se obavlja u kompresovanom formatu pomoću protokol bafera, preko HTTP
- Predviđa se prelazak na gRPC, kada se obezbedi da sve instance Prometheus ekosistema mogu da rade preko HTTP/2



<https://prometheus.io/docs/prometheus/latest/storage/>

SELF MONITORING

- prometheus.yml

global:

scrape_interval: 15s

evaluation_interval: 15s

scrape_configs:

- job_name: 'prometheus'

scrape_interval: 5s

static_configs:

- targets: ['localhost:9090']

- ./prometheus.exe --config.file=prometheus.yml

```
# HELP go_gc_duration_seconds A summary of the pause duration of garbage collection cycles.
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 0
go_gc_duration_seconds{quantile="0.25"} 0
go_gc_duration_seconds{quantile="0.5"} 0
go_gc_duration_seconds{quantile="0.75"} 0.000511
go_gc_duration_seconds{quantile="1"} 0.0100134
go_gc_duration_seconds_sum 0.0135646
go_gc_duration_seconds_count 12
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go_goroutines 41
# HELP go_info Information about the Go environment.
# TYPE go_info gauge
go_info{version="go1.16.2"} 1
# HELP go_memstats_alloc_bytes Number of bytes allocated and still in use.
# TYPE go_memstats_alloc_bytes gauge
go_memstats_alloc_bytes 3.9824512e+07
# HELP go_memstats_alloc_bytes_total Total number of bytes allocated, even if freed.
# TYPE go_memstats_alloc_bytes_total counter
go_memstats_alloc_bytes_total 3.63066808e+08
# HELP go_memstats_buck_hash_sys_bytes Number of bytes used by the profiling bucket hash table.
# TYPE go_memstats_buck_hash_sys_bytes gauge
go_memstats_buck_hash_sys_bytes 1.488154e+06
# HELP go_memstats_frees_total Total number of frees.
# TYPE go_memstats_frees_total counter
```


TARGETS

Targets

All Unhealthy Collapse All

node (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:9182/metrics	UP	instance="localhost:9182" job="node"	4.89s ago	617.872ms	

prometheus (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:9090/metrics	UP	instance="localhost:9090" job="prometheus"	573.000ms ago	4.731ms	

pushgateway (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:9091/metrics	UP	instance="localhost:9091" job="pushgateway"	1.561s ago	1.556ms	

python (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:8000/metrics	UP	instance="localhost:8000" job="python"	3.145s ago	317.075ms	

QUERIES

Prometheus Alerts Graph Status ▾ Help Classic UI

[Remove Panel](#)


🔍 job:process_cpu_seconds:rate5m

🔄 Execute

Table Graph

Load time: 20ms Resolution: 3s Result series: 1

15m + < End time > Res. (s) 📊 📈



job:process_cpu_seconds:rate5m (job="node")

2021-04-23 22:59:30 +00:00 23:01 23:02
job:process_cpu_seconds:rate5m: 0.05805124102536289
job: node

THIRD PARTY MONITORING

- Eksporter
- Windows exporter
https://github.com/prometheus-community/windows_exporter
- Metrike vezane za hardver i OS mašine
 - Iskorišćenost CPU, memorije, prostora na disku, upisi i čitanje sa diska, bandwidth mreže, trenutni broj aktivnih niti, ...

```
scrape_configs:  
  # The job name is added as a label 'job=<j>  
  - job_name: node  
    # metrics_path defaults to '/metrics'  
    # scheme defaults to 'http'.  
  
    # Override the global default and scrape  
    scrape_interval: 5s  
  
  static_configs:  
    - targets: ['localhost:9182']
```

SERVICE DISCOVERY

- Statički – znamo lokaciju izvora i ona je nepromenljiva

```
scrape_configs:  
  - job_name: 'node'  
    static_configs:  
      - targets: ['localhost:9182']
```

- Iz fajla – za service discovery sisteme koje Prometheus ne podržava out of the box

```
scrape_configs:  
  - job_name: file  
    file_sd_configs:  
      - files: '*.json'  
    [ {  
      "targets": [ "host1:9100", "host2:9100" ],  
      "labels": {  
        "team": "infra",  
        "job": "node"  
      }  
    },  
    {  
      "targets": [ "host1:9090" ],  
      "labels": {  
        "team": "monitoring",  
        "job": "prometheus"  
      }  
    }  
  ]
```

- Out of the box podrška za popularna SD rešenja

- Azure, Consul, DNS, EC2, OpenStack, File, Kubernetes, Marathon, Nerve, Serverset, and Triton
- Kubernetes

```
scrape_configs:  
  - job_name: 'k8services'  
    kubernetes_sd_configs:  
      - role: endpoints  
    relabel_configs:  
      - source_labels:  
        - __meta_kubernetes_namespace  
        - __meta_kubernetes_service_name  
        regex: default;kubernetes  
        action: drop  
      - source_labels:  
        - __meta_kubernetes_namespace  
        regex: default  
        action: keep  
      - source_labels: [__meta_kubernetes_service_name]  
        target_label: job
```

node, endpoints,
service, pod, ingress

METRIKE - INSTRUMENTACIJA

- Klijentske biblioteke – Go, Python, Java i Ruby
- Python - prometheus_client

```
import http.server
from prometheus_client import start_http_server


class MyHandler(http.server.BaseHTTPRequestHandler):
    def do_GET(self):
        self.send_response(200)
        self.end_headers()
        self.wfile.write(b"Hello World")

if __name__ == "__main__":
    start_http_server(8000)
    server = http.server.HTTPServer(('localhost', 8001), MyHandler)
    server.serve_forever()
```


```
# HELP python_gc_objects_collected_total Objects collected during gc
# TYPE python_gc_objects_collected_total counter
python_gc_objects_collected_total{generation="0"} 66.0
python_gc_objects_collected_total{generation="1"} 290.0
python_gc_objects_collected_total{generation="2"} 0.0
# HELP python_gc_objects_uncollectable_total Uncollectable object found during GC
# TYPE python_gc_objects_uncollectable_total counter
python_gc_objects_uncollectable_total{generation="0"} 0.0
python_gc_objects_uncollectable_total{generation="1"} 0.0
python_gc_objects_uncollectable_total{generation="2"} 0.0
# HELP python_gc_collections_total Number of times this generation was collected
# TYPE python_gc_collections_total counter
python_gc_collections_total{generation="0"} 47.0
python_gc_collections_total{generation="1"} 4.0
python_gc_collections_total{generation="2"} 0.0
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="8",patchlevel="7",version="3.8.7"} 1.0
# HELP hello_worlds_total Hello Worlds requested.
# TYPE hello_worlds_total counter
hello_worlds_total{method="GET",path="/"} 9.0
hello_worlds_total{method="GET",path="/favicon.ico"} 9.0
# HELP hello_worlds_created Hello Worlds requested.
# TYPE hello_worlds_created gauge
hello_worlds_created{method="GET",path="/"} 1.6192098147284164e+09
hello_worlds_created{method="GET",path="/favicon.ico"} 1.6192098147494187e+09
# HELP hello_world_exceptions_total Exceptions serving Hello World.
# TYPE hello_world_exceptions_total counter
hello_world_exceptions_total 5.0
# HELP hello_world_exceptions_created Exceptions serving Hello World.
# TYPE hello_world_exceptions_created gauge
hello_world_exceptions_created 1.6192098107888548e+09
# HELP hello_world_sales_euro_total Euros made serving Hello World.
# TYPE hello_world_sales_euro_total counter
hello_world_sales_euro_total 3.862290149433532
# HELP hello_world_sales_euro_created Euros made serving Hello World.
# TYPE hello_world_sales_euro_created gauge
hello_world_sales_euro_created 1.6192098107888548e+09
# HELP hello_worlds_inprogress Number of Hello Worlds in progress.
# TYPE hello_worlds_inprogress gauge
hello_worlds_inprogress 0.0
# HELP hello_world_last_time_seconds The last time a Hello World was served.
# TYPE hello_world_last_time_seconds gauge
hello_world_last_time_seconds 1.619209826976455e+09
```

METRIKE - INSTRUMENTACIJA


- Counter
- Gauge
- Summary
- Histogram



Histogram



Counter



Gauge

COUNTER

- Najčešće se koriste da bi merili broj izvršenja nekog dela koda

```
from prometheus_client import Counter
```

```
REQUESTS = Counter('hello_worlds_total', 'Hello Worlds requested.', labelnames=['path', 'method'])  
EXCEPTIONS = Counter('hello_world_exceptions_total', 'Exceptions serving Hello World.')  
SALES = Counter('hello_world_sales_euro_total', 'Euros made serving Hello World.')
```

```
REQUESTS.labels(self.path, self.command).inc()  
SALES.inc(euros)
```

```
with EXCEPTIONS.count_exceptions():  
    # Code which may raise an exception
```

```
@EXCEPTIONS.count_exceptions()  
def func():
```

- PromQL

rate(hello_world_exceptions_total[1m])

rate(hello_worlds_total[1m])

rate(hello_world_exceptions_total[1m]) /

sum without(method, path)(rate(hello_worlds_total[1m]))

GAUGE

- Trenutno stanje nekog resursa
- Npr. Broj aktivnih niti, zauzeće keš memorije, poslednji put kada je neki zahtev obrađen, ...
- Inc, dec i set

```
import time
from prometheus_client import Gauge

INPROGRESS = Gauge('hello_worlds_inprogress',
    'Number of Hello Worlds in progress.')
LAST = Gauge('hello_world_last_time_seconds',
    'The last time a Hello World was served.')

class MyHandler(http.server.BaseHTTPRequestHandler):
    def do_GET(self):
        INPROGRESS.inc()
        self.send_response(200)
        self.end_headers()
        self.wfile.write(b"Hello World")
        LAST.set(time.time())
        INPROGRESS.dec()
        LAST.set_to_current_time()
```

`time()` - `hello_world_last_time_seconds`

SUMMARY

- Pogodno za praćenje latencije
- `Observe()`
- Generiše 2 vremenske serije
 - `_count` – broj poziva observe funkcije
 - `_sum` – suma vrednosti prosleđenih observe funkciji

```
rate(hello_world_latency_seconds_sum[1m]) /  
rate(hello_world_latency_seconds_count[1m])
```

```
import time  
from prometheus_client import Summary  
  
LATENCY = Summary('hello_world_latency_seconds',  
                  'Time for a request Hello World.')
```

```
class MyHandler(http.server.BaseHTTPRequestHandler):  
    def do_GET(self):  
        start = time.time()  
        self.send_response(200)  
        self.end_headers()  
        self.wfile.write(b"Hello World")  
        LATENCY.observe(time.time() - start)
```

```
from prometheus_client import Summary  
  
LATENCY = Summary('hello_world_latency_seconds',  
                  'Time for a request Hello World.')
```

```
class MyHandler(http.server.BaseHTTPRequestHandler):  
    @LATENCY.time()  
    def do_GET(self):  
        self.send_response(200)  
        self.end_headers()  
        self.wfile.write(b"Hello World")
```

HISTOGRAM

- Quantiles and percentiles
- Observe() funkcija, isto kao za Summary
- Kreira `_count` i `_sum` vremenske serije (kao Summary) i niz vremenskih serija `_bucket`, gde je svaka od njih counter događaja koji upadaju u njen opseg

```
# HELP hello_world_latency_histogram_seconds Time for a request Hello World.
# TYPE hello_world_latency_histogram_seconds histogram
hello_world_latency_histogram_seconds_bucket{le="0.0002"} 27.0
hello_world_latency_histogram_seconds_bucket{le="0.0004"} 48.0
hello_world_latency_histogram_seconds_bucket{le="0.0008"} 50.0
hello_world_latency_histogram_seconds_bucket{le="0.0016"} 50.0
hello_world_latency_histogram_seconds_bucket{le="0.0032"} 50.0
hello_world_latency_histogram_seconds_bucket{le="0.0064"} 50.0
hello_world_latency_histogram_seconds_bucket{le="0.0128"} 50.0
hello_world_latency_histogram_seconds_bucket{le="0.0256"} 50.0
hello_world_latency_histogram_seconds_bucket{le="0.0512"} 50.0
hello_world_latency_histogram_seconds_bucket{le="+Inf"} 50.0
hello_world_latency_histogram_seconds_count 50.0
hello_world_latency_histogram_seconds_sum 0.010813600000204815
```

```
LATENCYH = Histogram('hello_world_latency_histogram_seconds',
                    'Time for a request Hello World.',
                    buckets=[0.0001 * 2**x for x in range(1, 10)])
```

```
histogram_quantile(0.95,
rate(hello_world_latency_histogram_seconds_bucket[1m]))
```

LABELE

- Daju dimenzionalnosti podacima
- Key/value parovi
- Vrednost labela i ime metrike zajedno jedinstveno identifikuju vremensku seriju podataka

```
http_requests_login_total  
http_requests_logout_total  
http_requests_adduser_total  
http_requests_comment_total  
http_requests_view_total
```

VS

```
http_requests_total{path="/login"}  
http_requests_total{path="/logout"}  
http_requests_total{path="/adduser"}  
http_requests_total{path="/comment"}  
http_requests_total{path="/view"}
```

- Target labela – identifikuju metu od koje Prometheus pribavlja podatke
 - Arhitektura, id aplikacije, id datacentra, production/development, tim, instanca
 - Prometheus ih dodaje u procesu prikupljanja podataka
- Labela za instrumentaciju – potiču iz procesa instrumentacije
 - Informacije koje su poznate unutar aplikacije/biblioteke – tip HTTP zahteva, baza kojoj se obraća
- PromQL ne pravi razliku između toga kog je tipa labela

LABEL PATTERNS

- Prometheus može da pamti samo 64-bitne float vrednosti, ne i stringove
- Labele su stringovi

Enum

- Stanje resursa – STARTING, RUNNING, STOPPING, TERMINATED
- Kodiranje brojevima nije pogodno za izračunavanja

```
# HELP gauge The current state of resources.
# TYPE gauge resource_state
resource_state{resource_state="STARTING",resource="blaa"} 0
resource_state{resource_state="RUNNING",resource="blaa"} 1
resource_state{resource_state="STOPPING",resource="blaa"} 0
resource_state{resource_state="TERMINATED",resource="blaa"} 0
```

avg_over_time(resource_state[1h])

without(resource)(resource_state)

```
from threading import Lock
from prometheus_client.core import GaugeMetricFamily, REGISTRY

class StateMetric(object):
    def __init__(self):
        self._resource_states = {}
        self._STATES = ["STARTING", "RUNNING", "STOPPING", "TERMINATED",]
        self._mutex = Lock()

    def set_state(self, resource, state):
        with self._mutex:
            self._resource_states[resource] = state

    def collect(self):
        family = GaugeMetricFamily("resource_state", "The current state of resources.",
                                    labels=["resource_state", "resource"])
        with self._mutex:
            for resource, state in self._resource_states.items():
                for s in self._STATES:
                    family.add_metric([s, resource], 1 if s == state else 0)
        yield family
```

LABEL PATTERNS

Info

- Za informacije koje nisu pogodne da budu target labels
- Gauge sa vrednošću 1 i labelama za sve željene informacije

```
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="5",patchlevel="2",
            version="3.5.2"} 1.0
```

```
up
* on (instance, job) group_left(version)
my_python_info
```

```
version_info = {
    "implementation": "CPython",
    "major": "3",
    "minor": "5",
    "patchlevel": "2",
    "version": "3.5.2"
}


INFO = Gauge("my_python_info",
             "Python platform information",
             labelnames=version_info.keys())

INFO.labels(**version_info).set(1)
```

NAZIVI METRIKA

DOBRA PRAKSA

- self-explanatory
- Hijerarhijski model imena – od najopštijeg ka specifičnijim opisima
- Prva reč – domen, namespace
- Osnovne jedinice
- Jedinica na kraju imena metrike
- 1 metrika = 1 značenje, labele određuju kontekst
- Svaka jedinstvena kombinacija vrednosti labela => zasebna vremenska serija
- Nazivi counter-a treba da se završavaju sa _total
- Izbegavati sufikse _count, _sum i _bucket jer se oni već koriste kod histograma i summary



UNIT TESTIRANJE INSTRUMENTACIJE

- Za testiranje nekih kritičnih metrika
- REGISTRY omogućava pristup raspoloživim podacima kroz kod
- get_sample_value funkcija

```
import unittest
from prometheus_client import Counter, REGISTRY

FOOS = Counter('foos_total', 'The number of foo calls.')

def foo():
    FOOS.inc()

class TestFoo(unittest.TestCase):
    def test_counter_inc(self):
        before = REGISTRY.get_sample_value('foos_total')
        foo()
        after = REGISTRY.get_sample_value('foos_total')
        self.assertEqual(1, after - before)
```

INSTRUMENTACIJA

- Online-serving sistemi – sinhroni interaktivni sistemi
 - Requests
 - Errors
 - Duration – latencija
- Offline-serving sistemi – pipeline za obradu podataka
 - Utilisation – zasićenost servisa
 - Saturation – količina posla na čekanju
 - Errors
- 10 miliona metrika i 1000 čvorova
- 1 nova metrika za svaku instancu je 0.01% resursa
- Histogrami uvode visoku kardinalnost
 - Za kardinalnost 100 – 1 nova metrika je 1% resursa
- Rule of thumb – 10 najvećih metrika zauzima polovinu resursa
 - Fokus na njihovoj optimizaciji

PRAVILA

- PromQL izrazi koji se rodovno evaluiraju
- Agregirani rezultati za kasnije korišćenje – ubrzava rad dashboard-ova
- Mala periodična obrada vs obrada velike količine podataka u trenutku izvršenja upita

```
global:  
  scrape_interval: 15s  
  evaluation_interval: 15s  
rule_files:  
  - 'prometheus.rules.yml'
```

```
groups:  
- name: cpu-node  
  rules:  
- record: job_instance_mode:process_cpu_seconds:avg_rate5m  
  expr: avg by (job, instance, mode) (rate(process_cpu_seconds_total[5m]))  
- record: job:process_cpu_seconds:rate5m  
  expr: sum without(instance) (rate(process_cpu_seconds_total{job="node"}[5m]))  
- record: job_device:windows_logical_disk_read_bytes:rate5m  
  expr: sum without(instance) (rate(windows_logical_disk_read_bytes_total{job="node"}[5m]))  
- record: job:node_disk_read_bytes:rate5m  
  expr: sum without(device) (job_device:windows_logical_disk_read_bytes:rate5m{job="node"})
```

PRAVILA

cpu-node			13.750s ago	1.235ms
Rule	State	Error	Last Evaluation	Evaluation Time
record: job_instance_mode:process_cpu_seconds:avg_rate5m expr: avg by(job, instance, mode) (rate(process_cpu_seconds_total[5m]))	OK		13.750s ago	0.141ms
record: job:process_cpu_seconds:rate5m expr: sum without(instance) (rate(process_cpu_seconds_total{job="node"}[5m]))	OK		13.750s ago	0.545ms
record: job_device:windows_logical_disk_read_bytes:rate5m expr: sum without(instance) (rate(windows_logical_disk_read_bytes_total{job="node"}[5m]))	OK		13.750s ago	0s
record: job:node_disk_read_bytes:rate5m expr: sum without(device) (job_device:windows_logical_disk_read_bytes:rate5m{job="node"})	OK		13.750s ago	0.549ms

ALERTS

- Specijalna pravila
- Pišu se u istom fajlu kao i pravila
- Umesto record, alert

```
- name: alerting
rules:
- alert: InstanceDown
  expr: up == 0
  for: 1m
- alert: ManyInstancesDown
  expr: >
    (
      avg without(instance) (up{job="node"}) < 0.5
      and on()
      hour() > 9 < 17
    )
```

alerting			6.962s ago	1.173ms
Rule	State	Error	Last Evaluation	Evaluation Time
alert: InstanceDown expr: up == 0 for: 1m	OK		6.964s ago	0.627ms
alert: ManyInstancesDown expr: (avg without(instance) (up{job="node"}) < 0.5 and on() hour() > 9 < 17)	OK		6.963s ago	0.546ms

ALERTS

▼ ManyInstancesDown (0 active)

name: `ManyInstancesDown`

expr: `(avg without(instance) (up{job="node"}) < 0.5 and on() hour() > 9 < 17)`

▼ InstanceDown (1 active)

name: `InstanceDown`

expr: `up == 0`

for: 1m

Labels	State	Active Since	Value
<code>alertname=InstanceDown</code> <code>instance=localhost:9091</code> <code>job=pushgateway</code>	PENDING	2021-04-23T23:26:21.009856531Z	0

▼ InstanceDown (1 active)

name: `InstanceDown`

expr: `up == 0`

for: 1m

Labels	State	Active Since	Value
<code>alertname=InstanceDown</code> <code>instance=localhost:9091</code> <code>job=pushgateway</code>	FIRING	2021-04-23T23:26:21.009856531Z	0

ALERTMANAGER

Alertmanager Alerts Silences Status Help

New Silence

Filter

Group

Receiver: All

☐ Silenced

☐ Inhibited

+

 Silence

Custom matcher, e.g. `env="production"`

+ Expand all groups

— alertname="InstanceDown"

+

1 alert

23:27:21, 2021-04-23 (UTC)

 Source

 Silence

instance="localhost:9091" +

job="pushgateway" +

monitor="codelab-monitor" +

ALERTMANAGER

- Inhibicija
 - Sprečavanje da se pošalje notifikacija za neki događaj u određenim uslovima
 - Ako je neki ozbiljniji alert u firing modu, ne okidati manje ozbiljne alertove
- Utišavanje
 - Privremeno onesposobljavanje alerta
 - Npr. Gašenje servisa zbog održavanja
- Rutiranje
 - Routing tree – putanje kojima se šalju različite notifikacije
 - Različiti timovi, production/development, ...
- Grupisanje
 - Za sve događaje iz istog reka/datacentra/servisa slati samo 1 notifikaciju
- Prigušivanje i ponavljanje
 - Sprečavanje spama - za slučaj da se pojavi novi alert iz grupe za koju je već poslata notifikacija, taj alert će da se priguši
 - Po potrebi će ista notifikacija da se pošalje više puta ako se ne vidi reakcija
- Obaveštavanje
 - Nakon svih prethodnih koraka, obaveštenja se šalju na prijemu adresu

BATCH POSLOVI

- Poslovi koji se obavljaju periodično, odrade neku obradu i nestanu


```
import random
from prometheus_client import CollectorRegistry, Gauge, pushadd_to_gateway


registry = CollectorRegistry()
duration = Gauge('my_job_duration_seconds', 'Duration of my batch job in seconds', registry=registry)

try:
    with duration.time():
        # Random job
        rand = random.random()
        if rand < 0.3:
            num = rand * 95
        else:
            num = rand * 85 / rand + 20
        rand = random.random()
        if rand < 0.2:
            raise Exception
    # This only runs if there wasn't an exception.
    g = Gauge('my_job_last_success_seconds', 'Last time my batch job successfully finished', registry=registry)
    g.set_to_current_time()
finally:
    pushadd_to_gateway('localhost:9091', job='batch', registry=registry)
```

PUSHGATEWAY

Pushgateway Metrics Status Help

 job="batch"


 my_job_duration_seconds Duration of my batch job in seconds GAUGE last pushed: 2021-04-23T22:11:52+02:00

Labels


Value

instance="" job="batch"

0.0000052999999999858716

 my_job_last_success_seconds Last time my batch job successfully finished GAUGE last pushed: 2021-04-23T22:11:52+02:00

 push_failure_time_seconds Last Unix time when changing this group in the Pushgateway failed. GAUGE last pushed: 2021-04-23T22:11:01+02:00

 push_time_seconds Last Unix time when changing this group in the Pushgateway succeeded. GAUGE last pushed: 2021-04-23T22:11:52+02:00

GRAFANA



LITERATURA

- Prometheus dokumentacija - <https://prometheus.io/>
- Brian Brazil – „Prometheus Up & Running – Infrastructure and Application performance monitoring“ – O'Reilly
- Google Cloud, Cloud Architecture Center
 - Remote monitoring and alerting for IoT - <https://cloud.google.com/architecture/remote-monitoring-and-alerting-for-iot>
 - Using Prometheus and Grafana for IoT monitoring - <https://cloud.google.com/community/tutorials/cloud-iot-prometheus-monitoring>
- Introduction to monitoring Microservices with Prometheus
<https://winderresearch.com/introduction-to-monitoring-microservices-with-prometheus/>